

In [13]: *# QUESTION 1 - IMPLEMENTATION OF LeNet Model using Non-Linear Activation*

```
In [1]: import cv2
import matplotlib.pyplot as plt
import tensorflow as tf
import os
from tensorflow.contrib.layers import flatten
from sklearn.utils import shuffle
import numpy as np
from sklearn.model_selection import train_test_split

base_path = os.path.abspath('')
dataset_path = os.path.join(base_path, 'PyCharmProjects/LeNet5_CNN/101_O
batch_size = 1024
n_epochs = 60
batches = 400
keep_prob = tf.placeholder(tf.float32)

def read_dataset(dataset_path):
    images = []
    labels = []
    label = 0
    ##Python3
    for paths in os.walk(dataset_path).__next__()[1]:
        dir = os.path.join(dataset_path, paths)
        walk = os.walk(dir).__next__() ##Python3
        for image in walk[2]:
            if image.endswith('.jpg'):
                images.append(os.path.join(dir, image))
                labels.append(label)
        label += 1

    imgs = []
    for img in images:
        imgs.append(cv2.imread(img, cv2.COLOR_BGR2RGB))
    return imgs, labels

imgs, labels = read_dataset(dataset_path)
xx = tf.placeholder(tf.float32, (None, 64, 64, 3))
yy = tf.placeholder(tf.int32, None)
one_hot_Y = tf.one_hot(yy, 101)

def LeNet(x, keep_prob):
    # Conv Layer 1
    weight1 = tf.get_variable("weight0", initializer=tf.truncated_normal
```

```

bias1 = tf.get_variable(name="b_b1", shape=[32],
                        initializer=tf.random_normal_initializer(stddev=0.01))
conv_layer1 = tf.nn.conv2d(x, weight1, strides=[1, 1, 1, 1], padding='SAME')
conv_layer1 = tf.nn.relu(conv_layer1)
pooling1 = tf.nn.max_pool(conv_layer1, ksize=[1, 4, 4, 1], strides=[1, 4, 4, 1],
                           padding='SAME')

# Conv layer 2
weight2 = tf.Variable(tf.truncated_normal(shape=[5, 5, 32, 64], stddev=0.01))
bias2 = tf.get_variable(name="b_b2", shape=[64],
                        initializer=tf.random_normal_initializer(stddev=0.01))
conv_layer2 = tf.nn.conv2d(pooling1, weight2, strides=[1, 1, 1, 1], padding='SAME')
conv_layer2 = tf.nn.relu(conv_layer2)
pooling2 = tf.nn.max_pool(conv_layer2, ksize=[1, 4, 4, 1], strides=[1, 4, 4, 1],
                           padding='SAME')

# Flatten layer
flatten_layer = flatten(pooling2)

# FC layer 1
fc1_weight1 = tf.get_variable("fc1_weight_1", initializer=tf.truncated_normal_initializer(stddev=0.01),
                              regularizer=tf.contrib.layers.l2_regularizer(0.001))
fc1_bias = tf.get_variable(name="fc_01_bias_", shape=[1024],
                           initializer=tf.random_normal_initializer(stddev=0.01))
fc1 = tf.matmul(flatten_layer, fc1_weight1) + fc1_bias
fc1 = tf.nn.relu(fc1)
# applied DropOut
# drop_out_1 = tf.nn.dropout(fc1, keep_prob)

# FC Layer 2
fc2_weight2 = tf.get_variable("fc_02_weight2_", initializer=tf.truncated_normal_initializer(stddev=0.01),
                              regularizer=tf.contrib.layers.l2_regularizer(0.001))
fc2_bias = tf.get_variable(name="fc_02_bias_", shape=[84],
                           initializer=tf.random_normal_initializer(stddev=0.01))
fc2 = tf.matmul(fc1, fc2_weight2) + fc2_bias
fc2 = tf.nn.relu(fc2)

# apply DropOut to hidden layer
# drop_out_2 = tf.nn.dropout(fc2, keep_prob) # DROP-OUT here

# FC Layer 3
fc3_weight3 = tf.Variable(tf.truncated_normal(shape=(84, 101), stddev=0.01))
fc3_bias = tf.get_variable(name="fc_03_bias_", shape=[101],
                           initializer=tf.random_normal_initializer(stddev=0.01))
logits = tf.matmul(fc2, fc3_weight3) + fc3_bias

return logits

###Training
with tf.variable_scope(tf.get_variable_scope(), reuse=tf.AUTO_REUSE):

```

```

logits = LeNet(xx, keep_prob)
cross_entropy = tf.nn.softmax_cross_entropy_with_logits_v2(logits=logits)
total_loss = tf.reduce_mean(cross_entropy) + tf.losses.get_regularization_loss()
LR = tf.train.AdamOptimizer(learning_rate=0.001).minimize(total_loss)

##Model Evaluation

pred = tf.equal(tf.argmax(logits, 1), tf.argmax(one_hot_Y, 1))
accu = tf.reduce_mean(tf.cast(pred, tf.float32))
saver = tf.train.Saver()

def visualize_weights(weights):

    plot_weight = np.moveaxis(weights, -1, 0)
    _, axs = plt.subplots(4, 8, figsize=(12, 12))
    axs = axs.flatten()
    for out, ax in zip(plot_weight, axs):
        v_min = plot_weight.min(axis=(0, 1), keepdims=True)
        v_max = plot_weight.max(axis=(0, 1), keepdims=True)
        out = (plot_weight - v_min) / (v_max - v_min)
        ax.imshow(plot_weight)
    plt.show()

def evaluate(X_data, y_data):
    num_examples = len(X_data)
    total_accuracy = 0
    sess = tf.get_default_session()
    for offset in range(0, num_examples, batches):
        batch_x, batch_y = X_data[offset:offset + batches], y_data[offset:offset + batches]
        accuracy = sess.run(accu, feed_dict={xx: batch_x, yy: batch_y, k: batch_y})
        total_accuracy += (accuracy * len(batch_x))
    return total_accuracy / num_examples

np_read_img = np.array(imgs)
##Normalize the numpy data
np_read_img = (np_read_img - np_read_img.min()) / (np.ptp(np_read_img))
np_labels = np.asarray(labels)

X_train, X_val, y_train, y_val = train_test_split(np_read_img, np_labels)
X_train, y_train = shuffle(X_train, y_train)

losses = []
with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    total_data = len(X_train)
    print("Training started...")
    for i in range(n_epochs):
        print("EPOCH {} ...".format(i + 1))

```

```

for offset in range(0, total_data, batches):
    end = offset + batches
    batch_x, batch_y = X_train[offset:end], y_train[offset:end]
    _, loss, acc = sess.run([LR, total_loss, accu], feed_dict={x
    print("Iter " + str(i) + ", Loss= " + \
          "{:.6f}".format(loss) + ", Training Accuracy= " + \
          "{:.5f}".format(acc))

validation_accuracy = evaluate(X_val, y_val)

print("Validation Accuracy = {:.3f}".format(validation_accuracy))
saver.save(sess, 'my-Lenet-model')
print("Model Saved")

```

```

Iter 58, Loss= 1.225723, Training Accuracy= 0.83750
Iter 58, Loss= 1.355819, Training Accuracy= 0.78250
Iter 58, Loss= 1.241440, Training Accuracy= 0.82500
Iter 58, Loss= 0.909884, Training Accuracy= 0.92344
Validation Accuracy = 0.571
Model Saved
EPOCH 60 ...
Iter 59, Loss= 1.346686, Training Accuracy= 0.78500
Iter 59, Loss= 1.316886, Training Accuracy= 0.79500
Iter 59, Loss= 1.401469, Training Accuracy= 0.77500
Iter 59, Loss= 1.370118, Training Accuracy= 0.77500
Iter 59, Loss= 1.276228, Training Accuracy= 0.80750
Iter 59, Loss= 1.309770, Training Accuracy= 0.79250
Iter 59, Loss= 1.266339, Training Accuracy= 0.83000
Iter 59, Loss= 1.200205, Training Accuracy= 0.82750
Iter 59, Loss= 1.217261, Training Accuracy= 0.82500
Iter 59, Loss= 1.363045, Training Accuracy= 0.78500
Iter 59, Loss= 1.380965, Training Accuracy= 0.78500
Iter 59, Loss= 1.349538, Training Accuracy= 0.80750
Iter 59, Loss= 1.245702, Training Accuracy= 0.81500

```

In [ ]: *#QUESTION 2.1 - Implementation of LeNet using LINEAR Activation Function*

```

In [3]: import cv2
import matplotlib.pyplot as plt
import tensorflow as tf
import os
from tensorflow.contrib.layers import flatten
from sklearn.utils import shuffle
import numpy as np
from sklearn.model_selection import train_test_split

base_path = os.path.abspath('')
dataset_path = os.path.join(base_path, 'PyCharmProjects/LeNet5_CNN/101_O
batch_size = 1024
n_epochs = 60
batches = 100

```

```

batches = 400
keep_prob = tf.placeholder(tf.float32)

def read_dataset(dataset_path):
    images = []
    labels = []
    label = 0
    ##Python3
    for paths in os.walk(dataset_path).__next__()[1]:
        dir = os.path.join(dataset_path, paths)
        walk = os.walk(dir).__next__() ##Python3
        for image in walk[2]:
            if image.endswith('.jpg'):
                images.append(os.path.join(dir, image))
                labels.append(label)
        label += 1

    imgs = []
    for img in images:
        imgs.append(cv2.imread(img, cv2.COLOR_BGR2RGB))
    return imgs, labels

imgs, labels = read_dataset(dataset_path)
xx = tf.placeholder(tf.float32, (None, 64, 64, 3))
yy = tf.placeholder(tf.int32, None)
one_hot_Y = tf.one_hot(yy, 101)

def LeNet(x, keep_prob):
    # Conv Layer 1
    weight1 = tf.Variable(tf.truncated_normal(shape=[5, 5, 3, 32], stddev=0.1),
                           name="w1",
                           initializer=tf.random_normal_initializer(stddev=0.1))
    bias1 = tf.get_variable(name="bb1", shape=[32],
                             initializer=tf.random_normal_initializer(stddev=0.1))
    conv_layer1 = tf.nn.conv2d(x, weight1, strides=[1, 1, 1, 1], padding='SAME')
    # conv_layer1 = tf.nn.relu(conv_layer1)
    pooling1 = tf.nn.max_pool(conv_layer1, ksize=[1, 4, 4, 1], strides=[1, 4, 4, 1],
                              padding='SAME')

    # Conv layer 2
    weight2 = tf.Variable(tf.truncated_normal(shape=[5, 5, 32, 64], stddev=0.1),
                           name="w2",
                           initializer=tf.random_normal_initializer(stddev=0.1))
    bias2 = tf.get_variable(name="bb2", shape=[64],
                             initializer=tf.random_normal_initializer(stddev=0.1))
    conv_layer2 = tf.nn.conv2d(pooling1, weight2, strides=[1, 1, 1, 1], padding='SAME')
    # conv_layer2 = tf.nn.relu(conv_layer2)
    pooling2 = tf.nn.max_pool(conv_layer2, ksize=[1, 4, 4, 1], strides=[1, 4, 4, 1],
                              padding='SAME')

    # Flatten layer
    flatten_layer = flatten(pooling2)

    # FC layer 1

```

```

fc1_weight1 = tf.get_variable("fc1_weight_1", initializer=tf.trunca
# regularizer=tf.contrib.layers.l2_regularizer(scale=0.1))
fc1_bias = tf.get_variable(name="fc_01_bias_", shape=[1024],
                           initializer=tf.random_normal_initializer(
fc1 = tf.matmul(flatten_layer, fc1_weight1) + fc1_bias
# fc1 = tf.nn.relu(fc1)
# applied Dropout
# drop_out_1 = tf.nn.dropout(fc1, keep_prob)

# FC Layer 2
fc2_weight2 = tf.get_variable("fc_02_weight2_", initializer=tf.trunc
# regularizer=tf.contrib.layers.l2_regularizer(scale=0.1))
fc2_bias = tf.get_variable(name="fc_02_bias_", shape=[84],
                           initializer=tf.random_normal_initializer(
fc2 = tf.matmul(fc1, fc2_weight2) + fc2_bias
# fc2 = tf.nn.relu(fc2)

# apply Dropout to hidden layer
# drop_out_2 = tf.nn.dropout(fc2, keep_prob) # DROP-OUT here

# FC Layer 3
fc3_weight3 = tf.Variable(tf.truncated_normal(shape=(84, 101), stdde
fc3_bias = tf.get_variable(name="fc_03_bias_", shape=[101],
                           initializer=tf.random_normal_initializer(
logits = tf.matmul(fc2, fc3_weight3) + fc3_bias

return logits

###Training
with tf.variable_scope(tf.get_variable_scope(), reuse=tf.AUTO_REUSE):
    logits = LeNet(xx, keep_prob)
cross_entropy = tf.nn.softmax_cross_entropy_with_logits_v2(logits=logits
total_loss = tf.reduce_mean(cross_entropy) + tf.losses.get_regularizatio
LR = tf.train.AdamOptimizer(learning_rate=0.001).minimize(total_loss)

##Model Evaluation

pred = tf.equal(tf.argmax(logits, 1), tf.argmax(one_hot_Y, 1))
accu = tf.reduce_mean(tf.cast(pred, tf.float32))
saver = tf.train.Saver()

def evaluate(X_data, y_data):
    num_examples = len(X_data)
    total_accuracy = 0
    sess = tf.get_default_session()
    for offset in range(0, num_examples, batches):

```

```

        batch_x, batch_y = X_data[offset:offset + batches], y_data[offset:offset + batches]
        accuracy = sess.run(accu, feed_dict={xx: batch_x, yy: batch_y, k: batch_size})
        total_accuracy += (accuracy * len(batch_x))
    return total_accuracy / num_examples

np_read_img = np.array(imgs)
##Normalize the numpy data
np_read_img = (np_read_img - np_read_img.min()) / (np.ptp(np_read_img))
np_labels = np.asarray(labels)

X_train, X_val, y_train, y_val = train_test_split(np_read_img, np_labels)
X_train, y_train = shuffle(X_train, y_train)

losses = []
with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    total_data = len(X_train)
    print("Training started...")
    for i in range(n_epochs):
        print("EPOCH {} ...".format(i + 1))
        for offset in range(0, total_data, batches):
            end = offset + batches
            batch_x, batch_y = X_train[offset:end], y_train[offset:end]
            _, loss, acc = sess.run([LR, total_loss, accu], feed_dict={x: batch_x, y: batch_y})
            print("Iter " + str(i) + ", Loss= " + \
                  "{:.6f}".format(loss) + ", Training Accuracy= " + \
                  "{:.5f}".format(acc))

        validation_accuracy = evaluate(X_val, y_val)

        print("Validation Accuracy = {:.3f}".format(validation_accuracy))
        print()

    print("Completed")

```

```

Iter 59, Loss= 1.000389, Training Accuracy= 0.90250
Iter 59, Loss= 1.108266, Training Accuracy= 0.88250
Iter 59, Loss= 1.076686, Training Accuracy= 0.88250
Iter 59, Loss= 1.084220, Training Accuracy= 0.91250
Iter 59, Loss= 1.075355, Training Accuracy= 0.89250
Iter 59, Loss= 1.051593, Training Accuracy= 0.90000
Iter 59, Loss= 1.008677, Training Accuracy= 0.90250
Iter 59, Loss= 0.986395, Training Accuracy= 0.92250
Iter 59, Loss= 1.066077, Training Accuracy= 0.88750
Iter 59, Loss= 1.004974, Training Accuracy= 0.89500
Iter 59, Loss= 1.022638, Training Accuracy= 0.91000
Iter 59, Loss= 1.142298, Training Accuracy= 0.86500
Iter 59, Loss= 0.928250, Training Accuracy= 0.95000
Iter 59, Loss= 1.015103, Training Accuracy= 0.93750

```

Iter 59, Loss= 0.989270, Training Accuracy= 0.92500  
 Iter 59, Loss= 0.810233, Training Accuracy= 0.97129  
 Validation Accuracy = 0.606

Completed

```
In [ ]: # QUESTION2.1
# What happens if we use a linear activation function in all convolution
# Compare training and validation loss (cross-entropy) after and before

# ANSWER
# I have removed all the relu activation function from all convolutional
# This is now a straight-forward linear activation model. Using the line
# network might result in loosing the classification capability.

# From my observation of the results I achieved, using a linear activation
# provided a minimal decrease in the overall performance (i.e) it means a
# accuracy than the non linear activation function.

# This however would be due to the addition of regularization and other
# However in this section, we received a better Training accuracy with
# compared to section 1.

# The validation loss is slightly lower while we are using this linear function
# I could achieve around 60% validation accuracy while running the code

# In order to improve the accuracy I have tried to increase the number of
# performance of the model.
```

```
In [ ]: ##### QUESTION 2.2 - IMPLEMENTATION OF ADDING ADDITIONAL LAYERS TO THE MODEL
```

```
In [20]: import cv2
import matplotlib.pyplot as plt
import tensorflow as tf
import os
from tensorflow.contrib.layers import flatten
from sklearn.utils import shuffle
import numpy as np
from sklearn.model_selection import train_test_split

base_path = os.path.abspath('')
dataset_path = os.path.join(base_path, 'PyCharmProjects/LeNet5_CNN/101_Object
batch_size = 1024
n_epochs = 60
batches = 400
keep_prob = tf.placeholder(tf.float32)

def read_dataset(dataset_path):
```



```

images = []
labels = []
label = 0
##Python3
for paths in os.walk(dataset_path).__next__()[1]:
    dir = os.path.join(dataset_path, paths)
    walk = os.walk(dir).__next__() ##Python3
    for image in walk[2]:
        if image.endswith('.jpg'):
            images.append(os.path.join(dir, image))
            labels.append(label)
    label += 1

imgs = []
for img in images:
    imgs.append(cv2.imread(img, cv2.COLOR_BGR2RGB))
return imgs, labels

imgs, labels = read_dataset(dataset_path)
xx = tf.placeholder(tf.float32, (None, 64, 64, 3))
yy = tf.placeholder(tf.int32, None)
one_hot_Y = tf.one_hot(yy, 101)

def LeNet(x, keep_prob):
    # Conv Layer 1
    weight1 = tf.Variable(tf.truncated_normal(shape=[5, 5, 3, 32], stddev=0.1))
    bias1 = tf.get_variable(name="bias_1", shape=[32],
                            initializer=tf.random_normal_initializer(stddev=0.1))
    conv_layer1 = tf.nn.conv2d(x, weight1, strides=[1, 1, 1, 1], padding='SAME')
    # conv_layer1 = tf.nn.relu(conv_layer1)
    pooling1 = tf.nn.max_pool(conv_layer1, ksize=[1, 4, 4, 1], strides=[1, 4, 4, 1],
                              padding='SAME')

    # Conv layer 2
    weight2 = tf.Variable(tf.truncated_normal(shape=[5, 5, 32, 48], stddev=0.1))
    bias2 = tf.get_variable(name="bias_2", shape=[48],
                            initializer=tf.random_normal_initializer(stddev=0.1))
    conv_layer2 = tf.nn.conv2d(pooling1, weight2, strides=[1, 1, 1, 1], padding='SAME')
    # conv_layer2 = tf.nn.relu(conv_layer2)
    pooling2 = tf.nn.max_pool(conv_layer2, ksize=[1, 1, 1, 1], strides=[1, 1, 1, 1],
                              padding='SAME')

    # Conv layer 3
    weight3 = tf.Variable(tf.truncated_normal(shape=[5, 5, 48, 64], stddev=0.1))
    bias3 = tf.get_variable(name="bias_3", shape=[64],
                            initializer=tf.random_normal_initializer(stddev=0.1))
    conv_layer3 = tf.nn.conv2d(pooling2, weight3, strides=[1, 1, 1, 1], padding='SAME')
    # conv_layer3 = tf.nn.relu(conv_layer3)
    pooling3 = tf.nn.max_pool(conv_layer3, ksize=[1, 3, 3, 1], strides=[1, 3, 3, 1],
                              padding='SAME')

```

```

# Flatten layer
flatten_layer = flatten(pooling3)

# FC layer 1
fc1_weight1 = tf.get_variable("fc1_weight_1", initializer=tf.truncated_normal_initializer(
# regularizer=tf.contrib.layers.l2_regularizer(scale=0.1))
fc1_bias = tf.get_variable(name="fc1_bias_01", shape=[1024],
                           initializer=tf.random_normal_initializer(
fc1 = tf.matmul(flatten_layer, fc1_weight1) + fc1_bias
# fc1 = tf.nn.relu(fc1)
# applied Dropout
# drop_out_1 = tf.nn.dropout(fc1, keep_prob)

# FC Layer 2
fc2_weight2 = tf.get_variable("fc2_weight_2", initializer=tf.truncated_normal_initializer(
# regularizer=tf.contrib.layers.l2_regularizer(scale=0.1))
fc2_bias = tf.get_variable(name="fc2_bias_02", shape=[84],
                           initializer=tf.random_normal_initializer(
fc2 = tf.matmul(fc1, fc2_weight2) + fc2_bias
# fc2 = tf.nn.relu(fc2)

# apply Dropout to hidden layer
# drop_out_2 = tf.nn.dropout(fc2, keep_prob) # DROP-OUT here

# FC Layer 3
fc3_weight3 = tf.Variable(tf.truncated_normal(shape=(84, 101), stddev=0.1))
fc3_bias = tf.get_variable(name="fc3_bias_03", shape=[101],
                           initializer=tf.random_normal_initializer(
logits = tf.matmul(fc2, fc3_weight3) + fc3_bias

return logits

###Training
with tf.variable_scope(tf.get_variable_scope(), reuse=tf.AUTO_REUSE):
    logits = LeNet(xx, keep_prob)
    cross_entropy = tf.nn.softmax_cross_entropy_with_logits_v2(logits=logits)
    total_loss = tf.reduce_mean(cross_entropy) + tf.losses.get_regularization_loss(logits=logits)
    LR = tf.train.AdamOptimizer(learning_rate=0.001).minimize(total_loss)

###Model Evaluation

pred = tf.equal(tf.argmax(logits, 1), tf.argmax(one_hot_Y, 1))
accu = tf.reduce_mean(tf.cast(pred, tf.float32))
saver = tf.train.Saver()

def evaluate(X_data, y_data):
    num examples = len(X_data)

```

```

total_accuracy = 0
sess = tf.get_default_session()
for offset in range(0, num_examples, batches):
    batch_x, batch_y = X_data[offset:offset + batches], y_data[offset:offset + batches]
    accuracy = sess.run(accu, feed_dict={xx: batch_x, yy: batch_y, k: batch_y})
    total_accuracy += (accuracy * len(batch_x))
return total_accuracy / num_examples

np_read_img = np.array(imgs)
##Normalize the numpy data
np_read_img = (np_read_img - np_read_img.min()) / (np.ptp(np_read_img))
np_labels = np.asarray(labels)

X_train, X_val, y_train, y_val = train_test_split(np_read_img, np_labels)
X_train, y_train = shuffle(X_train, y_train)

losses = []
with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    total_data = len(X_train)
    print("Training started...")
    for i in range(n_epochs):
        print("EPOCH {} ...".format(i + 1))
        for offset in range(0, total_data, batches):
            end = offset + batches
            batch_x, batch_y = X_train[offset:end], y_train[offset:end]
            _, loss, acc = sess.run([LR, total_loss, accu], feed_dict={x: batch_x, y: batch_y})
            print("Iter " + str(i) + ", Loss= " + \
                  "{:.6f}".format(loss) + ", Training Accuracy= " + \
                  "{:.5f}".format(acc))

        validation_accuracy = evaluate(X_val, y_val)

        print("Validation Accuracy = {:.3f}".format(validation_accuracy))
        print()

    print("Completed")

```

```

Iter 59, Loss= 0.009592, Training Accuracy= 1.00000
Iter 59, Loss= 0.007784, Training Accuracy= 1.00000
Iter 59, Loss= 0.008497, Training Accuracy= 1.00000
Iter 59, Loss= 0.017825, Training Accuracy= 0.99500
Iter 59, Loss= 0.007193, Training Accuracy= 1.00000
Iter 59, Loss= 0.014982, Training Accuracy= 0.99750
Iter 59, Loss= 0.016547, Training Accuracy= 0.99750
Iter 59, Loss= 0.009378, Training Accuracy= 1.00000
Iter 59, Loss= 0.009305, Training Accuracy= 1.00000
Iter 59, Loss= 0.009577, Training Accuracy= 1.00000
Iter 59, Loss= 0.009400, Training Accuracy= 1.00000

```

```

Iter 59, Loss= 0.008480, Training Accuracy= 1.00000
Iter 59, Loss= 0.008509, Training Accuracy= 1.00000
Iter 59, Loss= 0.008051, Training Accuracy= 1.00000
Iter 59, Loss= 0.007974, Training Accuracy= 1.00000
Iter 59, Loss= 0.007052, Training Accuracy= 1.00000
Iter 59, Loss= 0.002890, Training Accuracy= 1.00000
Validation Accuracy = 0.578

```

Completed

In [ ]: *#Question 2.2: Can we compensate for the effect of removing the non-linear layers?*

*# Answer: In this section the number of hidden convolutional layers are increased. This increases the depth of the model and gives enough strength to the model. From the observation of the results, the validation loss is much lesser. However the Validation accuracy seems to be lesser now compared to the previous.*

*# Thus adding the number of layers proved to be improved at the initial stage as the above. So we cannot assure that adding the new layer would compensate the activation function.*

*# So in my case, adding the layers instead of removing the non-linear activation we avoid overfitting and resolve other issues.*

*# In order to get a much better performance in accuracy, we can try rescaling and also try to fine tune the hyper parameters to achieve better improvement.*

In [ ]: *##Assignment\_2\_q\_3. Visualizing the weights*

```

In [4]: def visualize_weights(weights):
    plot_weight = np.moveaxis(weights, -1, 0)
    _, axs = plt.subplots(4, 8, figsize=(12, 12))
    axs = axs.flatten()
    for out, ax in zip(plot_weight, axs):
        v_min = out.min(axis=(0, 1), keepdims=True)
        v_max = out.max(axis=(0, 1), keepdims=True)
        out = (out - v_min) / (v_max - v_min)
        ax.imshow(out)
    print("visualized")
    plt.show()

sess = tf.Session()
new_saver = tf.train.import_meta_graph('my-Lenet-model.meta')
new_saver.restore(sess, tf.train.latest_checkpoint('./'))
all_vars = tf.get_collection(tf.GraphKeys.GLOBAL_VARIABLES)
var_name = 'weight0'
graph = tf.get_default_graph()

```





```
w2 = graph.get_tensor_by_name(var_name + ':0')
w2_saved = sess.run(w2) # print out tensor
visualize_weights(w2_saved)
print("Completed")
```

WARNING:tensorflow:From /opt/anaconda3/lib/python3.7/site-packages/tensorflow/python/training/saver.py:1266: checkpoint\_exists (from tensorflow.python.training.checkpoint\_management) is deprecated and will be removed in a future version.

Instructions for updating:

Use standard file APIs to check for files with this prefix.

INFO:tensorflow:Restoring parameters from ./my-Lenet-model  
visualized



Completed

In [ ]:

